

Algorithms and Computability

Lecture 5: Parallel Algorithms

SW6 spring 2025

Simonas Šaltenis



AALBORG UNIVERSITY
DENMARK

Parallel algorithms



- Goals of the lecture:
 - *to understand the model of **dynamic multithreading** (aka **fork-join parallelism**);*
 - *to understand **work**, **span**, and **parallelism** — the concepts necessary for the analysis of parallel algorithms;*
 - *to understand the main ideas used when parallelizing algorithms;*
 - *to understand and be able to analyze the **parallel merge sort** algorithm.*

Fibonacci Numbers




- $F(n) = F(n-1) + F(n-2)$
- $F(0) = 0, F(1) = 1$
 - 0, 1, 1, 2, 3, 5, 8, 13, 21, 34 ...

```
Fibonacci(n)
01 if n ≤ 1 then return n
02 else
03   x = Fibonacci(n-1)
04   y = Fibonacci(n-2)
05 return x + y
```

- Straightforward recursive procedure is slow!
- Why? How slow?
- Let's draw the recursion tree

Fibonacci Numbers



Fibonacci Numbers



- How many summations are there $W(n)$?
 - $W(n) = W(n - 1) + W(n - 2) + 1$
 - $W(n) \geq 2W(n - 2) + 1$ and $W(1) = W(0) = 0$
 - *Is this:*
 - ♦ A: Arithmetic series?
 - ♦ B: Geometric series?
 - Go to [Socrative](#) and vote.

Fibonacci Numbers (4)



- How many summations are there $W(n)$?
 - $W(n) = W(n - 1) + W(n - 2) + 1$
 - $W(n) \geq 2W(n - 2) + 1$ and $W(1) = W(0) = 0$
 - *Is this:*
 - ♦ A: Arithmetic series?
 - ♦ B: Geometric series?
 - Go to [Socratic](#) and vote.
 - Solving the recurrence we get
$$W(n) \geq 2^{(n+2)/2} - 1 \approx 1.4^{n+2}$$
 - Precisely $W(n) = \Theta(\varphi^n)$, where φ is the *golden ratio* $(1+\sqrt{5})/2 \approx 1.618$
- Running time is *exponential*.

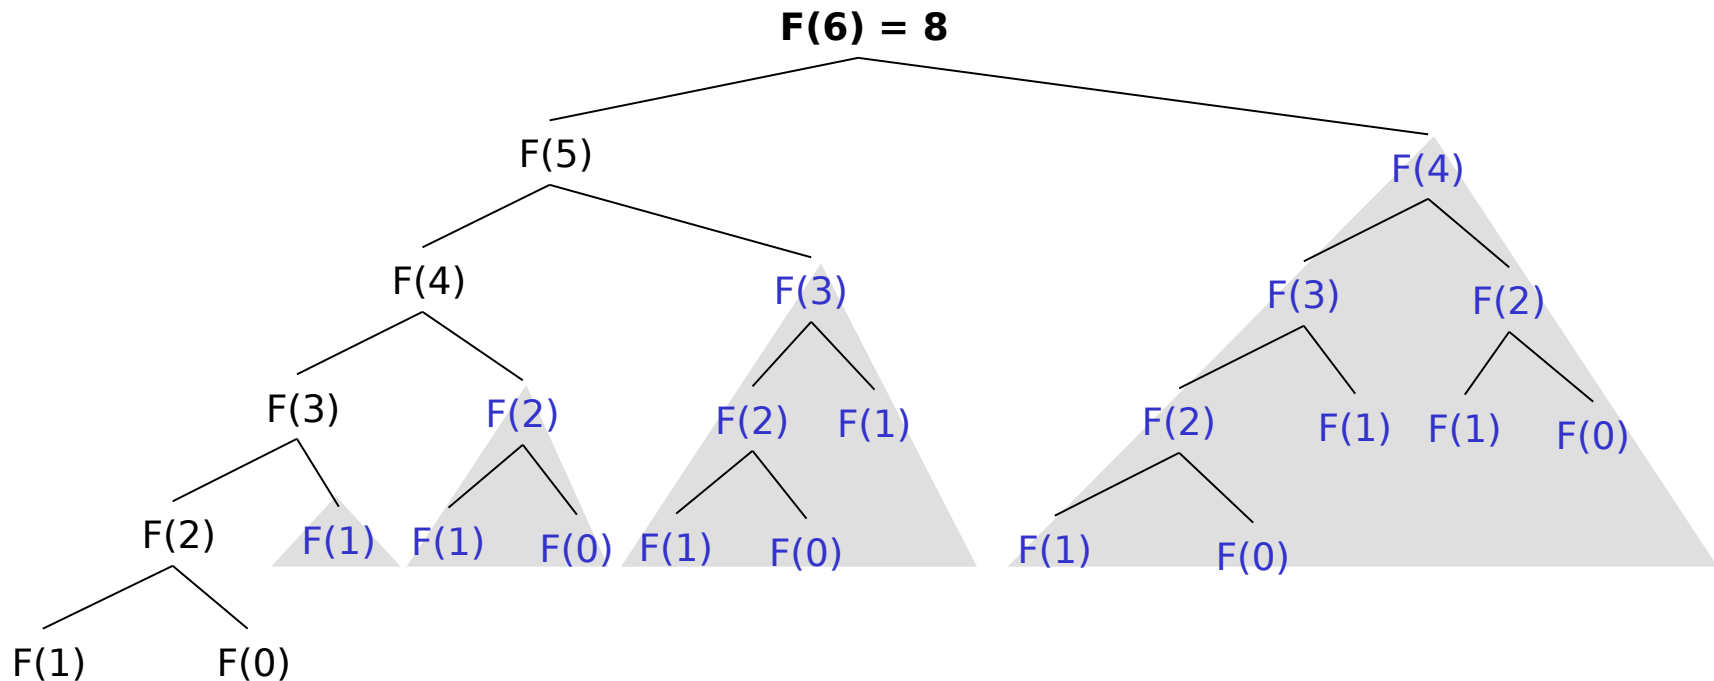
Multithreaded version



- What if we can do the two recursive calls in parallel
 - Using the so-called *nested parallelism*

```
FibonacciP(n)
01 if n ≤ 1 then return n
02 else
03   x = spawn FibonacciP(n-1)
04   y = FibonacciP(n-2)
05   sync
06 return x + y
```

FibonacciP analysis



- “Running time”:
 - $S(n) = \max(S(n - 1), S(n - 2)) + 1 = S(n - 1) + 1$
 - Thus $S(n) = \Theta(n)$

Work, Span, Parallelism



- Three main concepts (informally):
 - **Work**: the running time on a machine with one-processor (T_1).
 - ♦ Fibonacci: $\Theta(\varphi^n)$
 - **Span**: the running time on a machine with infinite processors (T_∞).
 - ♦ Fibonacci: $\Theta(n)$
 - **Parallelism** = $Work/Span$ – how many processors on average are used by the algorithm.
 - ♦ Fibonacci: $\Theta(\varphi^n / n)$
- More formally:
 - **Computation log / trace** – a DAG of serial strands of instructions (vertices) and dependencies (edges) between them.
 - Work = the number of vertices in the computation log.
 - Span = the length of the *longest path (critical path)* in the computation log.

Computation DAG



- Using an example of computing Fibonacci number of 4.

FibonacciP(n)

01 **if** $n \leq 1$ **then return** n

02 **else**

03 $x = \text{spawn FibonacciP}(n-1)$

04 $y = \text{FibonacciP}(n-2)$


05 **sync**

06 **return** $x + y$

● Lines 1-3

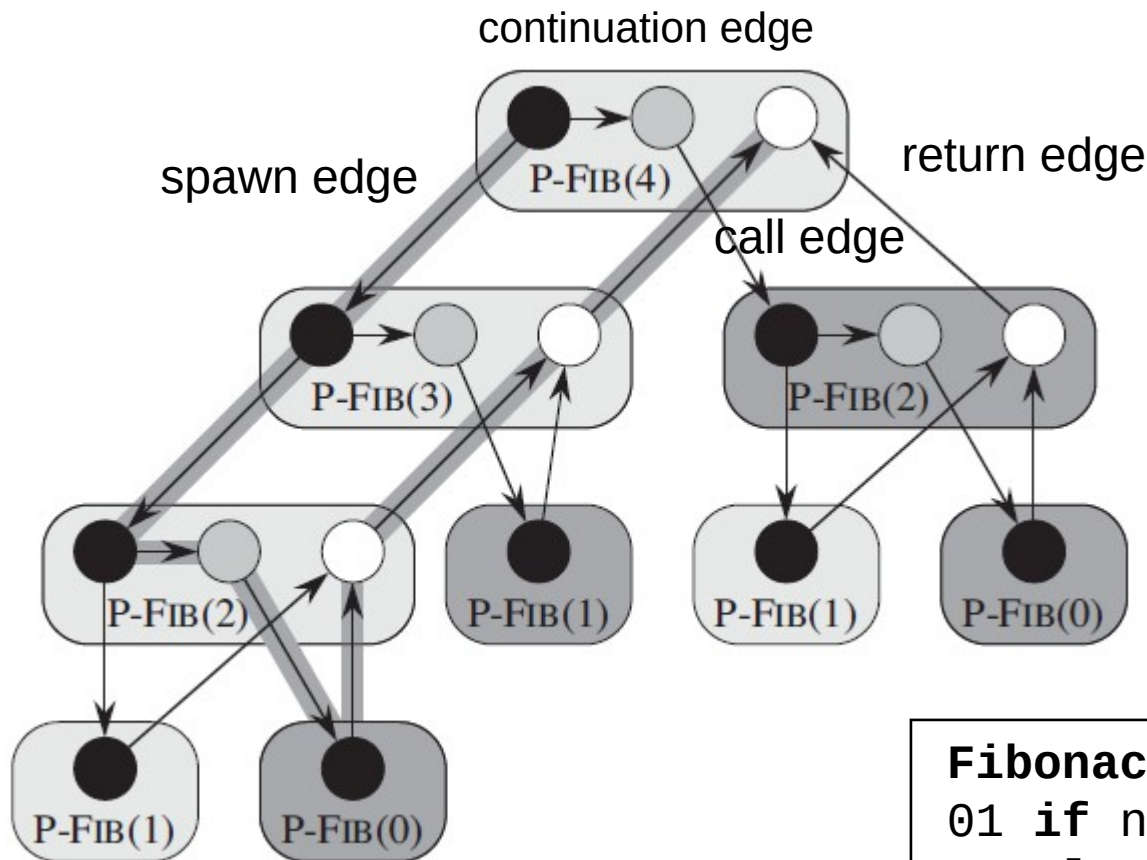
● Lines 4-5

○ Line 6




Computation DAG


Edge(u, v) means that
u must execute before v.



Spawned procedure



Called procedure



Work: number of vertices, 17
Span: the length of the longest path (critical path), 8

FibonacciP(n)

```

01 if n ≤ 1 then return n
02 else
03   x = spawn FibonacciP(n-1)
04   y = FibonacciP(n-2)
05   sync
06 return x + y
    
```

Work law and span law



- Notation
 - Work T_1
 - Span T_∞
 - Multithreaded computation on P processors: T_P
- Work law: $T_P \geq T_1 / P$
 - An ideal parallel computer with P processors can do at most P units of work.
- Span law: $T_P \geq T_\infty$
 - An ideal parallel computer with P processors cannot take less time than a machine with unlimited number of processors.

Assumptions



- The *fork-join parallelism (dynamic multithreading)* environment:
 - Shared-memory multi-core system
 - *Concurrency platform – task-parallel programming* :
 - ♦ Takes care of allocating work to physical threads (in other words: scheduling logical threads on physical threads)
 - ♦ Takes care of synchronization, consistent access to memory
 - Pseudocode keywords: *spawn*, *sync* and *parallel*
 - ♦ Indicates *potential* (or *logical*) parallelism: what *may* run in parallel.
 - We do not consider locking, race conditions, etc:
 - ♦ Parallel threads are *independent* – they work on separated items of data.
 - We abstract from actual physical scheduling:
 - ♦ It can be shown that simple greedy scheduling works well enough.

Speedup, Slackness



- When running on an actual system with P physical threads:
 - *Slackness* of a computation: $Parallelism / P$.
 - *What does it mean when slackness < 1 ? Slackness > 1 ?*
 - $Speedup = T_1 / T_P$.
 - *Perfect linear speedup*, when speedup = P .
 - Alternative way to think about **parallelism**:
 - ♦ *an upper bound on speedup*

Question



- Considering the case for computing P-Fib(4).
- We already know that the work $T_1 = 17$ and the span $T_\infty = 8$
- Consider the two setups $P = 2$ and $P = 3$, each setup corresponding to a machine with P processing units.
*Which setup is most likely to achieve the **perfect linear speedup**?*

To Summarize




Notation	Meaning
T_1	Work, the running time on a machine with one processor.
T_∞	Span, the running time on a machine with infinite processors.
T_P	The running time on a machine with P processors.
$T_P \geq T_1 / P$	Work law
$T_P \geq T_\infty$	Span law
T_1 / T_P	Speedup. Speedup must be $\leq P$ according to the work law. When speedup is equal to P, it achieves perfect speed up .
T_1 / T_∞	Parallelism. The maximum possible speedup that can be achieved on any number of processors
T_1 / PT_∞	Slackness = Parallelism/P. The larger the slackness, the more likely to achieve perfect speed up. When slackness is less than 1, it is impossible to achieve perfect speed up.

Computing span




- Sequential execution:

- Work and span: $T(A \text{ followed_by } B) = T(A) + T(B)$



- Parallel execution:

- Work: $T_1(A \text{ in_parallel_with } B) = T_1(A) + T_1(B)$
- Span: $T_\infty(A \text{ in_parallel_with } B) = \max(T_\infty(A), T_\infty(B))$



Goal of algorithm design



- Goal of the parallel algorithm design – *increase parallelism*.
 - Usually achieved by decreasing span (remember, parallelism = W/S)
 - It may pay off to slightly increase work, if span can be decreased significantly (in practice, relevant for highly parallel systems, such as supercomputers, but also GPUs)
- *What is this not?*
 - Task parallelism vs. data parallelism
 - Intra-operation parallelism vs. inter-operation parallelism.

Side Note: Efficient Fibonacci



- The efficient $O(n)$ serial algorithm:
 - Simple application of “*dynamic programming*” (or *memoized evaluation* of the recursive version)

```
FibonacciImproved(n)  
01 if  $n \leq 1$  then return  $n$   
02  $F_{im2} \leftarrow 0$   
03  $F_{im1} \leftarrow 1$   
04 for  $i \leftarrow 2$  to  $n$  do  
05    $F_i \leftarrow F_{im1} + F_{im2}$   
06    $F_{im2} \leftarrow F_{im1}$   
07    $F_{im1} \leftarrow F_i$   
05 return  $F_i$ 
```

- Can be actually done in $O(\lg n)$ additions and multiplications.

Parallel loops



- Denoted by the **parallel** keyword

```
ArrayCopy(A, B)
01 parallel for i = 1 to sizeof(A) do
02     B[i] = A[i]
```

- Analysis of span:
 - $S(n) = O(\lg n) + \max_i S_{\text{iteration}}(i)$
 - *Why?*
 - Parallel loop is implemented by divide-and-conquer

```
ArrayCopyRecursive(A, B, l, r)
01 if l > r then return
02 if l = r then B[l] = A[r]
03 else
04     q =  $\lceil (l+r)/2 \rceil$ 
05     spawn ArrayCopyRecursive(A, B, l, q-1)
06         ArrayCopyRecursive(A, B, q, r)
```

Examples



- Exchanging neighboring elements:

ArrayExchange(A)

```
01 parallel for i = 1 to |sizeof(A)/2| do
02     tmp      = A[i*2]
03     A[i*2]    = A[i*2-1]
04     A[i*2-1] = tmp
```



- Compute the largest stock price difference :

LargestSpike(A)

```
01 lspike = 0
02 parallel for i = 1 to sizeof(A)-1 do
03     if |A[i+1] - A[i]| > lspike then
04         lspike = |A[i+1] - A[i]|
05 return lspike
```



Examples



- Exchanging neighboring elements:

```
ArrayExchangeP(A)
01 parallel for i = 1 to |sizeof(A)/2| do
02   tmp[i] = A[i*2]
03   A[i*2] = A[i*2-1]
04   A[i*2-1] = tmp[i]
```

- Compute the largest stock price difference:

- Exercise

```
LargestSpikeP(A, l, r)
01 if trivial case then ...
02 else
03   divide...
04   ls = spawn LargestSpikeP(A, ...)
05   rs = LargestSpikeP(A, ...)
06   sync
07   combine
```

Parallelization techniques



- To parallelize an algorithm:
 - Convert *loops to parallel*, if possible
 - Rewrite parts of algorithm as *divide&conquer*
 - Brake the algorithm into more *phases* (that can then be parallelized)
 - Use *additional memory* to parallelize loops.

Parallelization techniques



- To parallelize an algorithm:
 - Convert *loops to parallel*, if possible
 - Rewrite parts of algorithm as *divide&conquer*
 - Brake the algorithm into more *phases* (that can then be parallelized)
 - Use *additional memory* to parallelize loops.
- In the video example of removing duplicates:
 - Nested loop
 - *inner loop as D&C*
 - *factor out a compaction phase to parallelize the main loop*
 - *make compaction into D&C*
 - *make the combine loop in compaction parallel by using additional memory...*

Merge Sort



```
Merge-Sort(A, p, r)
01 if p < r then
02     q = ⌊(p+r)/2⌋
03     Merge-Sort(A, p, q)
04     Merge-Sort(A, q+1, r)
05     Merge(A, p, q, r)
```

} *Divide*
 } *Conquer*
 } *Combine*

- Running time?

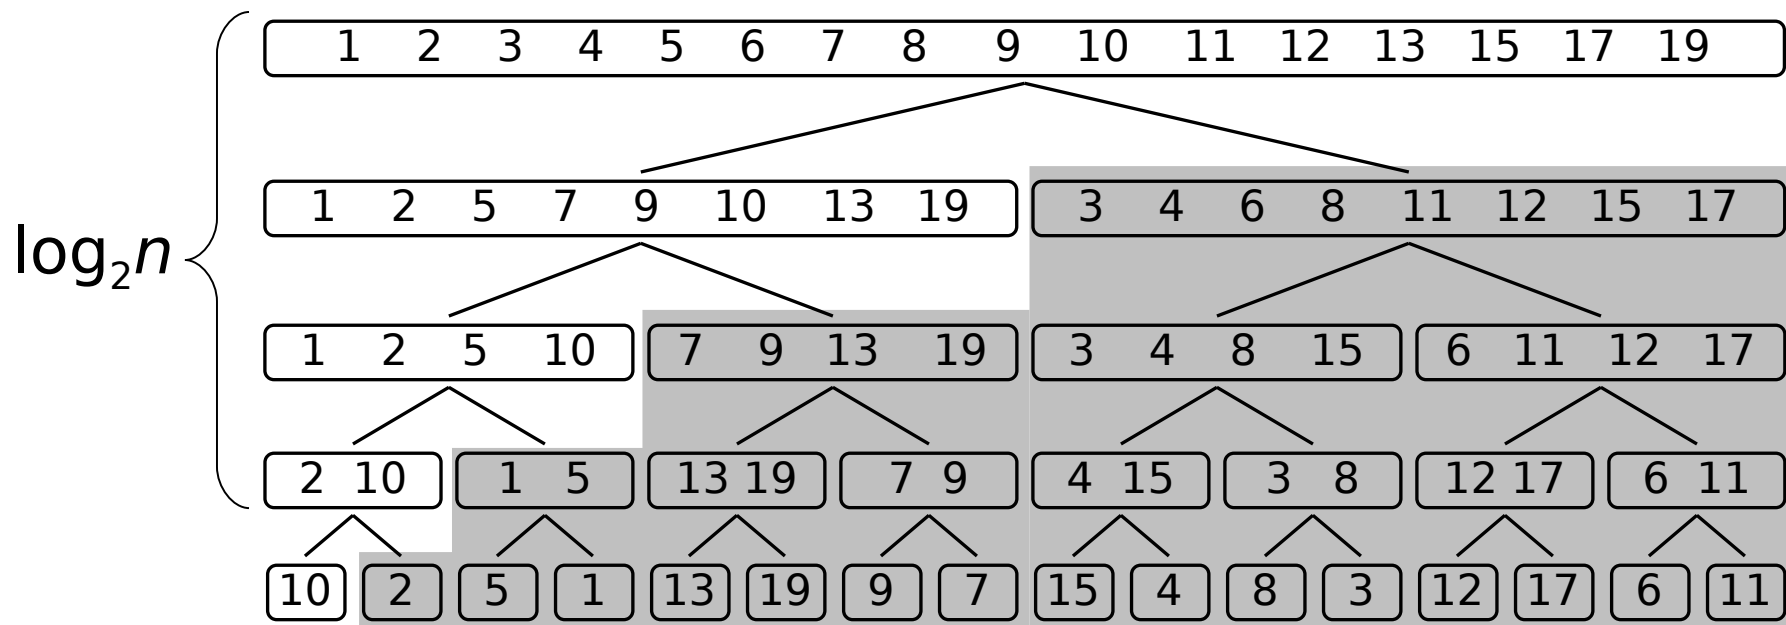
Parallelising Merge Sort



```
Merge-Sort'(A, p, r)
01 if p < r then
02   q = ⌊(p+r)/2⌋
03   spawn Merge-Sort'(A, p, q)
04   Merge-Sort'(A, q+1, r)
05   sync
06   Merge(A, p, q, r)
```

- Work
 - $W(n) = 2W(n/2) + \Theta(n)$
 - $W(n) = \Theta(n \lg n)$
- Span:
 - $S(n) = S(n/2) + \Theta(n)$
 - $S(n) = \Theta(n)$
- Parallelism: $W(n)/S(n) = \Theta(\lg n)$. Rather low...

Merge-Sort Recursion Tree

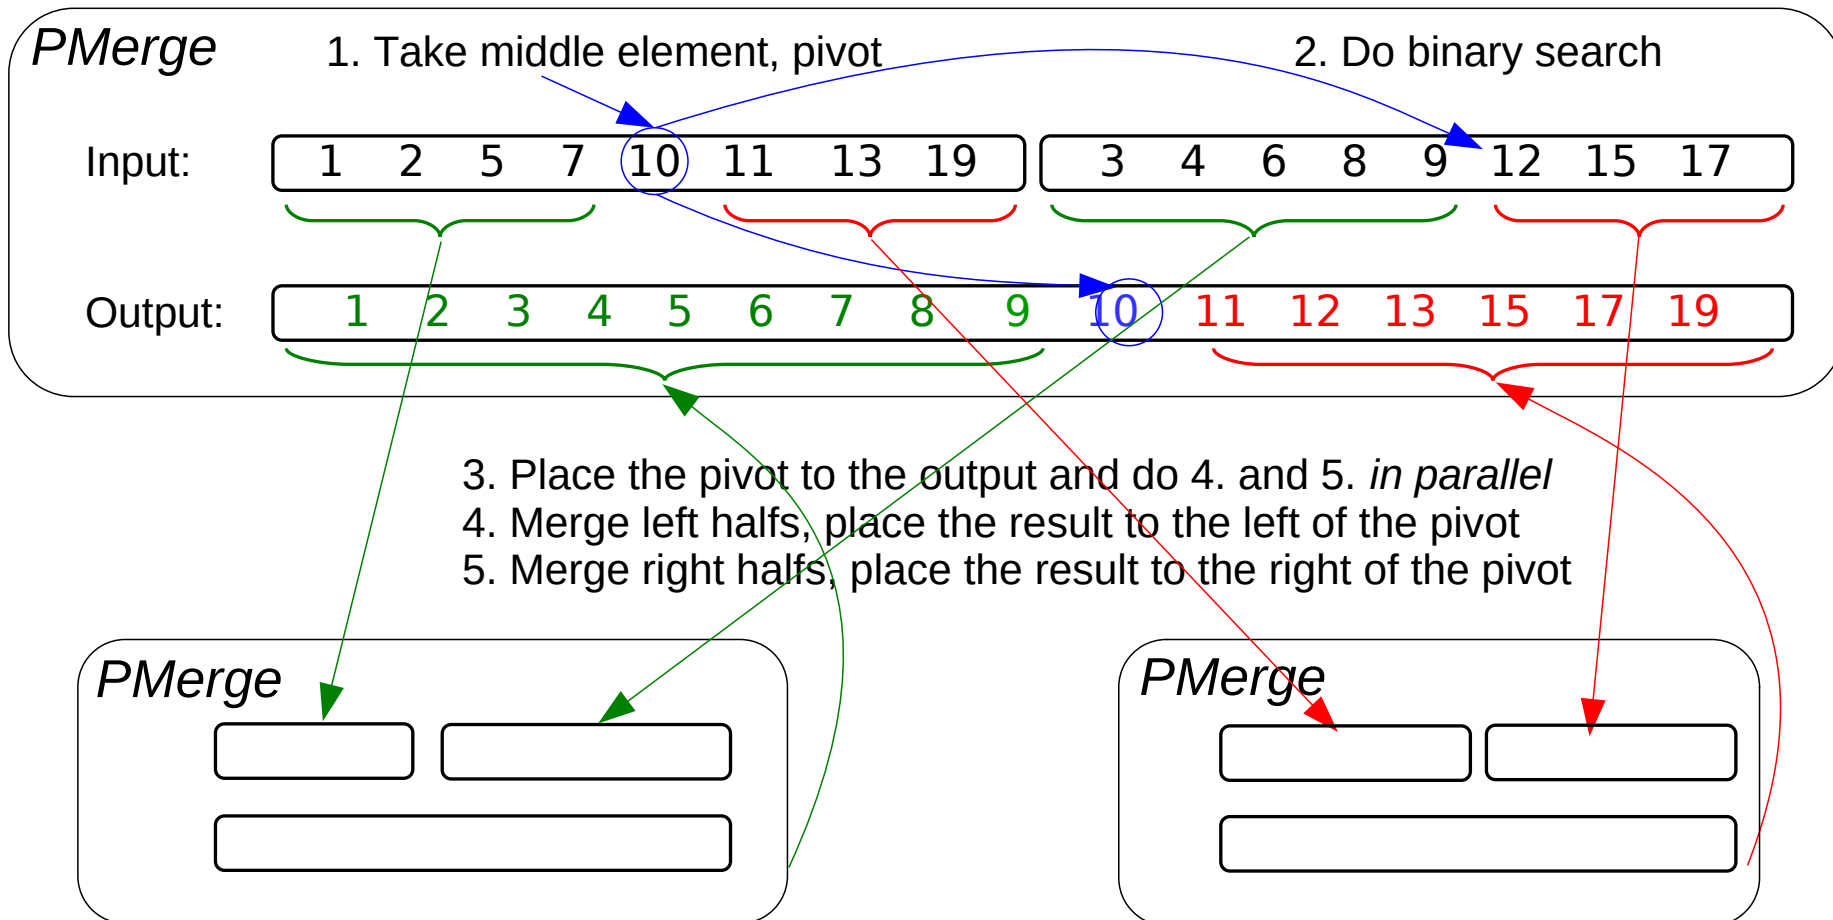


- Problem – merging is very serial:
 - At the top level, only one processor does $\Theta(n)$ work in serial!
 - At the second level, only two processors do $\Theta(n)$ work.
 - ...

Multithreaded merging



- Main idea – make the algorithm divide-and-conquer and use nested parallelism.



Multithreaded merging analysis



- One key idea: *do binary search in the smaller of the two arrays!*
 - This ensures that the *largest* of the two recursive calls works with *at most* $3n/4$ elements, where n = the sum of sizes of the two arrays.
 - *Why?*
- Span:
 - $S(n) = S(3n/4) + \Theta(\lg n)$
 - *What is the solution?*
 - $S(n) = \Theta(\lg^2 n)$
- Work:
 - Can be shown to be $\Theta(n)$.

Multithreaded merge sort



```
PMerge-Sort(A, p, r)
01 if p < r then
02     q = ⌊(p+r)/2⌋
03     spawn PMerge-Sort(A, p, q)
04     PMerge-Sort(A, q+1, r)
05     sync
06     PMerge(A, p, q, r)
```

- Work:
 - The same recurrence and solution: $W(n) = \Theta(n \lg n)$
- Span:
 - $S(n) = S(n/2) + \Theta(\lg^2 n)$
 - $S(n) = \Theta(\lg^3 n)$
- Parallelism:
 - $W(n) / S(n) = \Theta(n \lg n) / \Theta(\lg^3 n) = \Theta(n / \lg^2 n)$

To summarize



- Work: $\Theta(n \lg n)$

	Span of Merge	Total span	Parallelism
Naïve merge	$\Theta(n)$	$\Theta(n)$	$\Theta(\lg n)$
P-Merge	$\Theta(\lg^2 n)$	$\Theta(\lg^3 n)$	$\Theta(n / \lg^2 n)$

Goal of the parallel algorithm design



- Goal of the multi-threaded algorithm design – *increase parallelism*.
 - $Parallelism = work / span$
 - Usually achieved by decreasing span
 - ♦ MergeSort without P-Merge and with P-Merge
 - ♦ $\Theta(n)$ vs. $\Theta(\lg^3 n)$
 - It may pay off to slightly increase work, if span can be decreased significantly (in practice, relevant for highly parallel systems, such as supercomputers, GPUs)